Data-Driven Typesetting of Chinese Character Decompositions with `zitex`

Wen G. Gong
wen.gong.research@gmail.com

Abstract

Early Chinese writing, including the oracle bone script, contains many pictographs, and a character can be described by its sound, its shape and its meaning. To streamline the typesetting process of a document with chinese characters of pinyin, glosses and pictures, we create `zitex`, an extension for XeLaTeX that generates this layout from structured character data records. A character is a short entry giving its decomposition and the position of each part in the eleven-position Zi-Matrix notation. A Python program validates the data and writes a self-contained snippet per character and view: a one-line formula, a typeset equation with pinyin and glosses, a nested-box matrix, or a tree. A further view adds a row of pictures under the meanings (an oracle bone glyph, an image or an emoji), so that each part shows its sound, its meaning and a picture in rows that line up. A short LaTeX package includes those snippets. In a handout of 59 equations built from a database, correcting one gloss meant editing one line of data, and regenerating the 59 affected figures took under a second. The code is released under the MIT licence at `https://github.com/digital-duck/zinets-latex`.

1 Introduction

Early Chinese writing contains many pictographic signs. The oracle bone script of the late Shang dynasty, more than three thousand years ago, is the earliest substantial body of Chinese writing that survives, and the glyphs chosen in Section 4 for 日 (sun) and 月 (moon) look like a circle and a crescent. Modern characters differ greatly in form from such glyphs, and showing an ancient glyph next to the modern character is one way to present a character.

A character can be described along three dimensions: a sound (pinyin), a shape (the glyph, and how it is built from parts) and a meaning. Material about characters may present all three for each part of a character and for the character itself. In the author's experience this means manual layout: pinyin and glosses are placed above and below the characters, pictures are pasted and aligned, and the decomposition is drawn in a graphics program or in TikZ. The problem we address is a typesetting problem: to lay out sound, shape and meaning in a consistent way, generated from records and not positioned by hand. The records themselves (the decompositions, the glosses and the choice of glyphs) are still curated by people: `zitex` automates the layout, not the analysis.

Existing tools cover parts of this. LaTeX packages such as `xpinyin` add pinyin to Chinese text automatically [6], ideographic description sequences record how characters are composed [2], and datasets of oracle bone characters provide glyph images [3, 5]. We are not aware of a tool that lays these out together in one aligned figure, but we did not survey the field systematically (Section 6).

A decomposition can be written as an equation in which the parts combine into the character, for example 亻 + 弗 → **佛** (fó, Buddha). This is a representation of the chosen decomposition, used here for layout, and not a claim about how a character arose. The decompositions and the position codes used here follow [1].

This paper describes `zitex`, which typesets decompositions directly from data. Its aims are: (i) a figure is generated from a record, so that revising the record revises the figure; (ii) the notation gives a named position to every part, on top of the arrangement that an ideographic description sequence records, including levels with three or more parts; (iii) sound, shape and meaning are laid out the same way for every character, in rows that line up, with a picture that is an oracle bone glyph, an image or an emoji; (iv) the LaTeX side stays small, and the data and layout logic live in a program that can be tested; and (v) the source database is read only, unless the user asks otherwise (`db-build --in-place`). Section 4 shows glyphs for 日, 月, 明, 木 and 林 set next to their sounds and meanings.

2 Notation

A character is described level by level. At each level the parts have positions from a fixed set of eleven codes.

Code	Region
`L`, `R`, `U`, `D`	left, right, up, down
`LU`, `LD`, `RU`, `RD`	the four corners
`M`	the centre of gravity of the level
`Mi`, `Mo`	inside and outside (enclosing)

Table 1: The Zi-Matrix position codes, introduced in [1].

The *pattern* of a level is its position codes in reading order, such as `L-M` for 佛 or `LU-RU-M-LD-RD` for 器. A pattern determines the matching Unicode ideographic description operator where one exists: for instance `L-M` and `L-R` map to the left–right operator, and the enclosing patterns `Mo-D`, `Mo-U`, `Mo-R` and `Mo-M` map to the surrounding operators. Some arrangements have no single operator, for example four mouths around a dog in 器. For these `zitex` adds a few operators (such as ∴ ∵ ⁝ ⋮) and a general form that names the pattern. A *radical rule* resolves the ambiguity of two-part levels in which exactly one part is a Kangxi radical: a radical is peripheral, so it is placed at `L` (or `R`, when it stands on the right) and the other part at `M`, as in 佛 (亻 and 弗). When both parts are radicals, as in 明 and 林, the rule does not apply and the level is `L-R`. Positions are stored on every level, so a part that is itself decomposed (澡 inside 藻) carries its own pattern.

3 Design

Python does the parsing and layout, and LaTeX does the typesetting (Figure 1). `zitex extract` selects characters from a database and writes `chars.yaml`, `lexicon.yaml`, a `main.tex`, and every snippet the document asks for. The only interface between the two halves is the file name of a snippet, `u<hex>-<char>-<style>[-<lang>][@reading].tex`, where `u<hex>` is the code point of the character in hex (so the files sort by code point) and `<char>` is the character itself, and the package looks the file up and inputs it. A missing file gives a warning and a red mark, so a long build is not stopped. A language or picture-set name must not contain a hyphen, which separates the parts of the name, or a plus sign, which joins several picture sets (`pict=obs,pict=seal` is the style `pictmath-obs+seal`).

Installing it takes one command, `pip install zitex`, from the Python Package Index (it needs Python 3.10 or later and, for typesetting, XeLaTeX). The command `zitex setup` then writes a working folder with `zitex.sty`, a font configuration detected from the installed fonts, and the fallback

files, and `zitex extract -c 器, 藻 -o work/` fills it with a document that `xelatex main.tex` builds, so a first page needs no manual configuration.

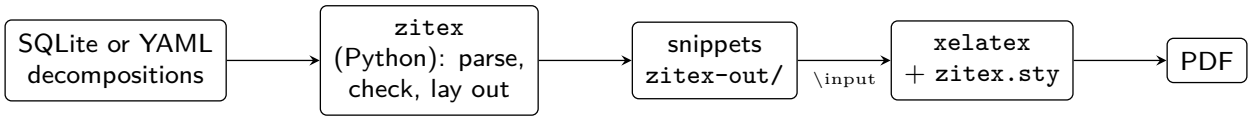


Figure 1: The workflow. Python writes one self-contained snippet per character and style, and XeLaTeX includes them.

Views. The equation view uses only `\text`, `\overset`, `\underbrace` and `\xrightarrow`, so it also works in MathJax. The matrix and the tree are TikZ pictures with coordinates computed in Python, which allows patterns of any size without a TikZ tree library.

Pictures. A character has sound, shape and meaning, and the equation view shows only the first and the last. The picture view (`\zimath[pict]{字}`) lays the equation out as one grid: for every part, the pinyin, the character, a short meaning and a picture, in rows that line up across the parts, so that the pictures share one line and the meanings another. A row that no part fills is left out, and a part without a picture has an empty cell. A picture is an image file, or else an emoji, or else nothing. XeLaTeX cannot print colour emoji, so an emoji is drawn into a PNG when the snippets are built, so the picture depends on the installed emoji font; the PNGs of this paper come from Noto Color Emoji and are in the source. Every picture is set in a box of the same size, which is a setting and not part of the data. The meaning is the gloss, or else the concept, or else the concept category, or else nothing. Pictures come in named sets (`\zimath[pict=obs]{字}`), so one document can show the same equations with emoji and with oracle bone glyphs. A picture that was asked for and cannot be made (a missing file, a word that is neither a file nor an emoji, no emoji font) is reported when the snippets are built, together with the parts that have no picture; `--strict` turns these reports into an error.

Readings and roles. Which part carries the meaning and which the sound is an interpretation. A character may store several named *readings*, selected per call with `reading=`, so that two interpretations can be set in one document without duplicating the data. With `--arrow role` the arrow of an equation is labelled with the part that is recorded as giving the sound.

Validation. `zitex check` reports unknown positions, patterns an operator does not allow, overlapping regions, and a radical placed at the centre. These are warnings and not errors, because the data being typeset is itself under review.

Fonts. A component is resolved in a fixed order: an explicit override, the primary font, fallback fonts, an image, and finally a boxed code point with a warning. Components that have no Unicode code point can therefore be set like any other.

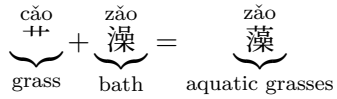
Data. Decompositions come from YAML files or from a SQLite database. Every command opens the database read only, except the opt-in `db-build --in-place`. `zitex` derives its own tables from it, and a hand revision is stored together with what it replaced, so that it can be reverted and is re-applied when the tables are rebuilt. A Streamlit application reviews and revises decompositions, and generates a PDF for a list of characters.

4 Example

The commands below make a self-contained folder and typeset three characters.

```
zitex setup   -o work/
zitex extract -c 器,藻,佛 -o work/
cd work && xelatex main.tex
```

The document asks for views with `\zimath{藻}`, `\zimatrix{器}`, `\zitree{藻}` and `\ziformula{佛}`. They give:



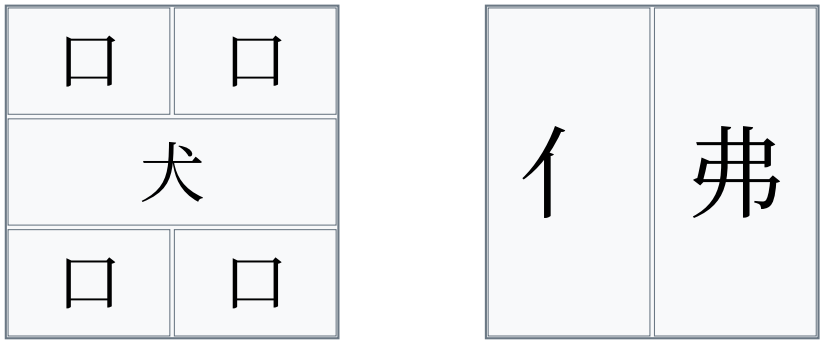


Figure 2: Zi-Matrices of 器 (four 口 at the corners and 犬 in the centre) and 佛 (亻 and 弗).

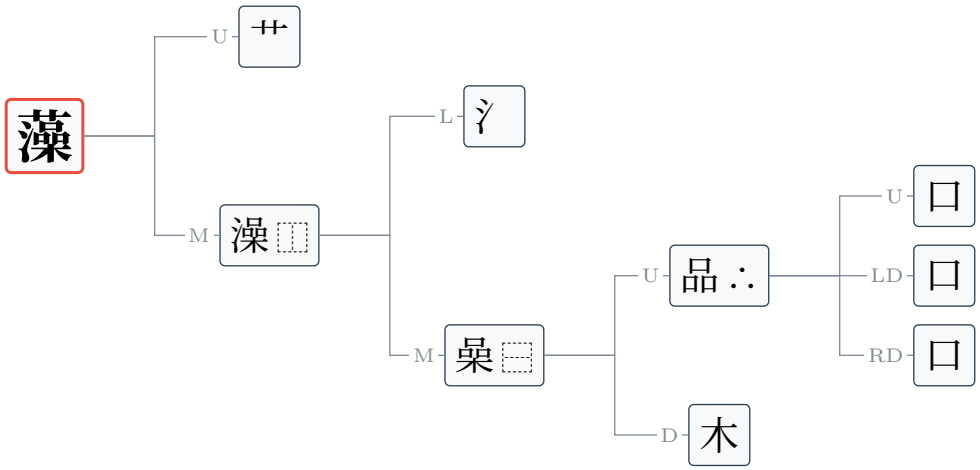


Figure 3: The decomposition tree of 藻. Leaves are parts that are not decomposed further.

The matrix keeps the geometry (Figure 2), and the tree shows the depth (Figure 3): 藻 has four levels, because 澡 and 桑 are decomposed too. Inline, `\ziformula{器}` gives 口 + 口 + 犬 + 口 + 口 → **器** (qì, device).

4.1 Sound, shape and meaning

The picture view adds the third dimension of a character. Each part carries its pinyin above, its meaning below, and a picture under the meaning. The input is a lexicon entry per character (here the pictures are emoji), the decomposition of the result, and the call in the document:

```
# lexicon.yaml
# (pict is the emoji character itself, or an image file;
# <emoji: sun> stands for it here)
日: {pinyin: rì,   gloss: {en: sun,    zh: 太阳}, pict: <emoji: sun>}
月: {pinyin: yuè,  gloss: {en: moon,   zh: 月亮}, pict: <emoji: moon>}
明: {pinyin: míng, gloss: {en: bright, zh: 明亮}, pict: <emoji: sparkles>}

\zimath[pict]{明}            % English meanings
\zimath[pict,lang=zh]{明}    % Chinese meanings
```

The figures below juxtapose modern decompositions, dictionary glosses and selected oracle bone glyph images. These are separate records, and placing them side by side does not establish a historical relationship between the modern components and the ancient forms.

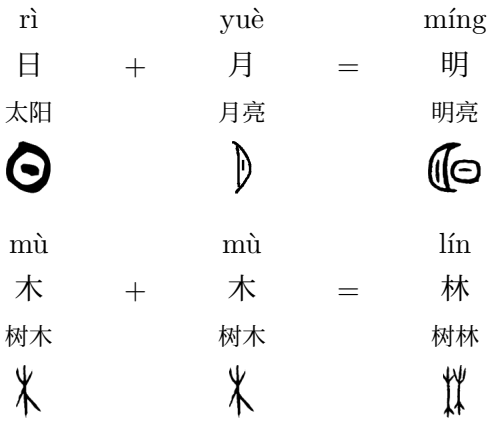
Emoji as the picture.



The meanings are in English in the first and third equation and in Chinese in the second: the language is a key of the call, and any language that the lexicon has can be used. The equations show a chosen decomposition and the glosses of its parts. They do not say why a part is there, and they are not a sum of the glosses: whether a part contributes meaning or sound is a matter of analysis on which sources differ, and a character can store several named readings of it (Section 3).

Oracle bone glyphs as the picture. The picture row can also show an ancient form of each part. The glyphs come from the HUST-OBC dataset of oracle bone characters [3], which is licensed CC BY-NC 4.0, one chosen glyph per character (Table 2). They are stored as `u<hex>.png` files and added to the lexicon as a second set of pictures, `pict_obs`. A call names the set it wants:

```
zitex render -i chars.yaml --scan main.tex \
    -x lexicon.yaml -x obs/lexicon.yaml -l en -l zh
\zimath[pict=obs,lang=zh]{明}
\zimath[pict=obs,lang=zh]{林}
```

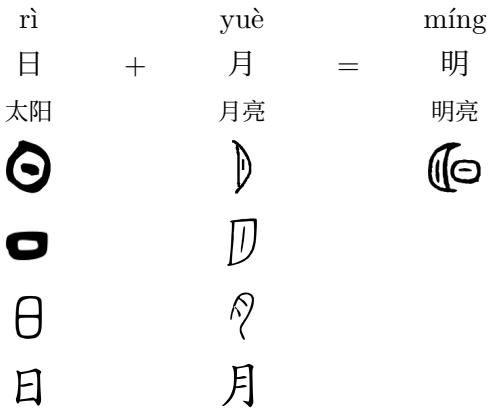


Character	HUST-OBC image (file name)	Kind of source
日	G_0773_前1.22.2合36169黃組	rubbing, Huangzu group
月	H_0815_605E1_12	handwritten copy (HWOBC)
明	G_0824_乙6672合11498賓組	rubbing, Binzu group
木	G_0829_乙4309合6477賓組	rubbing, Binzu group
林	G_0849_京都2308合34544歷組	rubbing, Lizu group

Table 2: The oracle bone glyph images shown above, from HUST-OBC [3] (CC BY-NC 4.0), one per character. The file name is the dataset's, without the extension. The last column is inferred from the file name. Each glyph is one sample, chosen by the author, and not a canonical form; one equation can mix glyphs of different periods and kinds.

Several forms in one equation. A picture row is one set, and a call can name more than one: each `pict=` key adds a row, in the order written, so the glyph of each part can be followed down the page from the oracle bone form to the modern one. The bronze, seal and kaiti sets come from hanziyuan.net [4], added to the lexicon in the same way as `pict_obs` (here `pict_bronze`, `pict_seal` and `pict_kaiti`). The rows line up under the same characters, and a part with no picture in a set leaves that cell empty.

```
\zimath[pict=obs,pict=bronze,pict=seal,pict=kaiti,lang=zh]{明}
```



The sets are independent: `\zimath[pict=obs,pict=kaiti]{明}` shows just the oracle bone and kaiti rows. The pictures are samples, one per form, and not canonical glyphs. Like the HUST-OBC images, they are shown only to demonstrate the layout and are not redistributed with the code.
The pictures differ from the emoji in what they show. The glyph chosen for 明 shows a crescent beside a small box (we do not identify the box), and the glyph chosen for 林 shows two tree-like signs side by side, as the modern decomposition 木 + 木 does. The picture row displays these records next to the modern decomposition and does not interpret either. The agreement for 林 and the disagreement for 明 show that the ancient form of a whole character need not contain the ancient forms of its modern parts. The gloss of each part is that of the modern character and may not be what the ancient glyph depicted.

4.2 Cost of a correction

The script `benchmarks/handout.py` builds a handout of equations, `\zimath{X}` for the characters of the database that contain a chosen component (口), then changes one line of `lexicon.yaml`, the gloss of that component, and rebuilds. For 59 equations the first build took 0.08 s for `zitex render --scan` and 0.41 s for one XeLaTeX run, and the rebuild after the edit took the same time and changed 59 of 59 snippets (one Linux machine with an Intel i7-13700F; three runs gave the same figures to 0.02 s). The same correction made by hand touches every equation that carries the gloss as text. The benchmark measures the cost of a correction and not the cost of writing the data, which remains with the author, and it uses only the equation view, which needs no pictures.

5 Use cases

We see three uses. The first is in daily use by the author; the other two are plans, and we have not measured them.

Typesetting a book. A book about Chinese characters repeats the same figures: the equation of a character, its tree, its matrix. With `zitex` each is one command, `\zimath{X}`, that reads the data and nothing else, so a correction to a gloss or a position is made once, in `chars.yaml` or `lexicon.yaml`, and every figure that shows it is rebuilt (Section 4, cost of a correction). The snippet folder is built locally and uploaded with the source, so the book compiles with plain XeLaTeX. The picture row takes an image per part, which a book can fill with emoji, oracle bone, bronze, seal or kaiti glyphs, as in Section 4. We have typeset only short documents so far, and not a full book.

Education. A teacher or a learner can make a handout that shows how a character is built from parts, with pinyin above each part and a gloss below it, and with the older forms of the parts in rows under them. The language of the gloss is a setting (`lang=zh` or `lang=fr`, besides English), and the example in this paper changes language without any change to the figure. A set of pictures can be swapped for another by naming it (`pict=kaiti` for learners of the modern form, `pict=obs` for the oldest). We have not tested handouts with students, so we make no claim about learning.

A dataset for vision models. For each part, the tool joins a decomposition (the operator and a named position for each part, in the database), a reading, a gloss, and an image of the part in several historical forms. Each image has a recorded source (the HUST-OBC file or the hanziyuan.net form), and every snippet is listed in a manifest with its character and style. This is the structure that a dataset of image and structure pairs for a vision-language model would need: for example, an oracle bone glyph paired with the decomposition and the modern form it is later written as. We have not built or released such a dataset, and we have not trained a model on one.

6 Related work and limitations

Related work. The Unicode Standard defines ideographic description characters, and community projects distribute ideographic description sequences for large character sets [2]. A sequence encodes the arrangement of the parts with operators (for example left–right), but not a named position for each part, and it does not produce a figure; `zitex` reads this form and adds positions per level. The package `xpinyin` annotates Chinese text with pinyin automatically [6]; `zitex` instead places pinyin, a gloss and a picture for each part of a decomposition, in one aligned figure. Oracle bone datasets such as HUST-OBC [3] and EVOBC, which concerns the evolution of oracle bone characters [5], provide the glyph images that the picture view can display; `zitex` does not classify or interpret them. Figures are produced with TikZ [8], and Chinese text with `xeCJK` [9], which `zitex` builds on. We did not survey other tools for drawing character decompositions, and a comparison is future work.

Implementation limits. The tool needs XeLaTeX and has not been tried with other engines. The picture view is not available in MathJax, because it uses images. A long equation with long meanings can be wider than the page: the width of a meaning and the number of senses shown are settings, and nothing scales the equation. Right-to-left text has not been tried. Snippets are produced by a Python program, which an online LaTeX service does not run at compile time, so the snippet folder is built locally and uploaded with the document. The snippet file names are the code point and the character (`u85fb-藻-math-fr.tex`), so they are not ASCII; they work on Linux, and we have not tried Windows or other file systems.

Data limits. The decompositions, glosses and the choice of glyphs are curated by hand and are not validated by the tool. The decomposition database used for development is a first iteration under review: about 220 rows of about 6,200 cannot be converted yet, and warnings remain on some characters. Glosses are first senses from a dictionary and are not curated. The oracle bone glyphs are a handful chosen by the author for these examples, and most elemental characters have no chosen glyph yet. We make no claim about the etymology of any character: the figures show the data.

Acknowledgements

This work began as a conversation with Google Gemini, and the building, testing and writing phases were assisted by Claude Code of Anthropic. The character decompositions, solution design, and result validation are the author's.

The English glosses are derived from CC-CEDICT (CC BY-SA 4.0) [7] and are attributed in the data. The oracle bone glyph images in Section 4 are from HUST-OBC [3], which is licensed CC BY-NC 4.0: they may be reused for non-commercial research with attribution, and are not covered by the MIT licence of the code. The bronze, seal and kaiti images are from hanziyuan.net [4], which states no licence (the rights stay with the site and its own sources); they are shown for demonstration only. The font subset in the source is cut from Noto Serif CJK SC (SIL Open Font License 1.1). The repository's `NOTICE` lists all third-party material and its licences.

References

[1] Wen G. Gong. *A New Exploration into Chinese Characters: from Simplification to Deeper Understanding*. arXiv:2502.19428, 2025.

[2] *cjkvi-ids: Ideographic Description Sequences for CJKV characters*. `https://github.com/cjkvi/cjkvi-ids`.

[3] Pengjie Wang et al. *An open dataset for oracle bone script recognition and decipherment*. arXiv:2401.15365, 2024.

[4] *Hanzi Yuan (hanziyuan.net): the origin and evolution of Chinese characters*. `https://hanziyuan.net`.

[5] Haisu Guan et al. *An open dataset for the evolution of oracle bone characters: EVOBC*. arXiv:2401.12467, 2024.

[6] *The xpinyin package*. CTAN, `https://ctan.org/pkg/xpinyin`.

[7] *CC-CEDICT: a Chinese-English dictionary*. `https://cc-cedict.org`.

[8] Till Tantau. *The TikZ and PGF Packages*. CTAN, `https://ctan.org/pkg/pgf`.

[9] *The xeCJK package*. CTAN, `https://ctan.org/pkg/xecjk`.